

In general, Apache Pig works on top of Hadoop. It is an analytical tool that analyzes large datasets that exist in the **Hadoop File System**. To analyze data using Apache Pig, we have to initially load the data into Apache Pig. This chapter explains how to load data to Apache Pig from HDFS.

Preparing HDFS

In MapReduce mode, Pig reads (loads) data from HDFS and stores the results back in HDFS. Therefore, let us start HDFS and create the following sample data in HDFS.

Student ID	First Name	Last Name	Phone	City
001	Rajiv	Reddy	9848022337	Hyderabad
002	siddarth	Battacharya	9848022338	Kolkata
003	Rajesh	Khanna	9848022339	Delhi
004	Preethi	Agarwal	9848022330	Pune
005	Trupthi	Mohanthi	9848022336	Bhuwaneshwar
006	Archana	Mishra	9848022335	Chennai

The above dataset contains personal details like id, first name, last name, phone number and city, of six students.

The input file of Pig contains each tuple/record in individual lines. And the entities of the record are separated by a delimiter (In our example we used “,”).

In the local file system, create an input file **student_data.txt** containing data as shown below.

```
001,Rajiv,Reddy,9848022337,Hyderabad
002,siddarth,Battacharya,9848022338,Kolkata
003,Rajesh,Khanna,9848022339,Delhi
004,Preethi,Agarwal,9848022330,Pune
005,Trupthi,Mohanthi,9848022336,Bhuwaneshwar
006,Archana,Mishra,9848022335,Chennai.
```

Now, move the file from the local file system to HDFS.

Verify whether the file has been moved into the HDFS.

You can load data into Apache Pig from the file system (HDFS/ Local) using **LOAD** operator of **Pig Latin**.

Syntax

The load statement consists of two parts divided by the “=” operator. On the left-hand side, we need to mention the name of the relation **where** we want to store the data, and on the right-hand side, we have to define **how** we store the data. Given below is the syntax of the **Load** operator.

```
Relation_name = LOAD 'Input file path' USING function as schema;
```

Where,

- **relation_name** – We have to mention the relation in which we want to store the data.
- **Input file path** – We have to mention the HDFS directory where the file is stored. (In MapReduce mode)
- **function** – We have to choose a function from the set of load functions provided by Apache Pig (**BinStorage**, **JsonLoader**, **PigStorage**, **TextLoader**).
- **Schema** – We have to define the schema of the data. We can define the required schema as follows –

```
(column1 : data type, column2 : data type, column3 : data type);
```

Note – We load the data without specifying the schema. In that case, the columns will be addressed as \$01, \$02, etc... (check).

Example

As an example, let us load the data in **student_data.txt** in Pig under the schema named **Student** using the **LOAD** command.

```
grunt> student = LOAD
'hdfs://localhost:9000/pig_data/student_data.txt'
  USING PigStorage(',')
  as ( id:int, firstname:chararray, lastname:chararray,
phone:chararray,
  city:chararray );
```

Following is the description of the above statement.

Relation name	We have stored the data in the schema student .					
Input file path	We are reading data from the file student_data.txt , which is in the /pig_data/ directory of HDFS.					
Storage function	We have used the PigStorage() function. It loads and stores data as structured text files. It takes a delimiter using which each entity of a tuple is separated, as a parameter. By default, it takes ‘\t’ as a parameter.					
schema	We have stored the data using the following schema.					
	column	id	firstname	lastname	phone	city
	datatype	int	char array	char array	char array	char array

Note – The **load** statement will simply load the data into the specified relation in Pig.

Dump Operator

The **Dump** operator is used to run the Pig Latin statements and display the results on the screen. It is generally used for debugging Purpose.

Syntax

Given below is the syntax of the **Dump** operator.

```
grunt> Dump Relation_Name
```

Now, let us print the contents of the relation using the **Dump operator** as shown below.

```
grunt> Dump student
```

Once you execute the above **Pig Latin** statement, it will start a MapReduce job to read data from HDFS.

Describe Operator

The **describe** operator is used to view the schema of a relation.

Syntax

The syntax of the **describe** operator is as follows –

```
grunt> Describe Relation_name
```

let us describe the relation named **student** and verify the schema as shown below.

```
grunt> describe student;
```

Output

Once you execute the above **Pig Latin** statement, it will produce the following output.

```
grunt> student: { id: int,firstname: chararray,lastname:
chararray,phone: chararray,city: chararray }
```

Explain Operator

The **explain** operator is used to display the logical, physical, and MapReduce execution plans of a relation.

Syntax

Given below is the syntax of the **explain** operator.

```
grunt> explain Relation_name;
```

let us explain the relation named student using the **explain** operator as shown below.

```
grunt> explain student;
```

Output

It will produce the following output.

\$ explain student;

```
2015-10-05 11:32:43,660 [main]
2015-10-05 11:32:43,660 [main] INFO
org.apache.pig.newplan.logical.optimizer
.LogicalPlanOptimizer -
{RULES_ENABLED=[AddForEach, ColumnMapKeyPrune,
ConstantCalculator,
GroupByConstParallelSetter, LimitOptimizer, LoadTypeCastInserter,
MergeFilter,
MergeForEach, PartitionFilterOptimizer,
PredicatePushdownOptimizer,
PushDownForEachFlatten, PushUpFilter, SplitFilter,
StreamTypeCastInserter]}
#-----
# New Logical Plan:
#-----
student: (Name: LOStore Schema:
id#31:int,firstname#32:chararray,lastname#33:chararray,phone#34:c
hararray,city#
35:chararray)
|
|---student: (Name: LOForEach Schema:
id#31:int,firstname#32:chararray,lastname#33:chararray,phone#34:c
hararray,city#
35:chararray)
|      |
|      | (Name: LOGenerate[false,false,false,false,false] Schema:
id#31:int,firstname#32:chararray,lastname#33:chararray,phone#34:c
hararray,city#
35:chararray)ColumnPrune:InputUids=[34, 35, 32, 33,
31]ColumnPrune:OutputUids=[34, 35, 32, 33, 31]
|      |      |
|      |      | (Name: Cast Type: int Uid: 31)
|      |      | |---id:(Name: Project Type: bytearray
Uid: 31 Input: 0 Column: (*))
|      |      |
|      |      | (Name: Cast Type: chararray Uid: 32)
|      |      |
|      |      | |---firstname:(Name: Project Type: bytearray Uid: 32
Input: 1
Column: (*))
|      |      |
|      |      | (Name: Cast Type: chararray Uid: 33)
|      |      |
|      |      | |---lastname:(Name: Project Type: bytearray Uid: 33
Input: 2
Column: (*))
|      |      |
|      |      | (Name: Cast Type: chararray Uid: 34)
|      |      |
|      |      | |---phone:(Name: Project Type: bytearray Uid: 34
Input: 3 Column:
(*))
|      |      |
|      |      | (Name: Cast Type: chararray Uid: 35)
```

```

|      |      |
|      |      |---city:(Name: Project Type: bytearray Uid: 35
Input: 4 Column:
(*)
|      |
|      |---(Name: LOInnerLoad[0] Schema: id#31:bytearray)
|      |
|      |---(Name: LOInnerLoad[1] Schema: firstname#32:bytearray)
|      |
|      |---(Name: LOInnerLoad[2] Schema: lastname#33:bytearray)
|      |
|      |---(Name: LOInnerLoad[3] Schema: phone#34:bytearray)
|      |
|      |---(Name: LOInnerLoad[4] Schema: city#35:bytearray)
|
|---student: (Name: LOLoad Schema:
id#31:bytearray,firstname#32:bytearray,lastname#33:bytearray,phon
e#34:bytearray
,city#35:bytearray)RequiredFields:null
#-----
# Physical Plan: #-----
student: Store(fakefile:org.apache.pig.builtin.PigStorage) -
scope-36
|
|---student: New For Each(false,false,false,false,false) [bag] -
scope-35
|      |
|      |Cast[int] - scope-21
|      |
|      |---Project[bytearray][0] - scope-20
|      |
|      |Cast[chararray] - scope-24
|      |
|      |---Project[bytearray][1] - scope-23
|      |
|      |Cast[chararray] - scope-27
|      |
|      |---Project[bytearray][2] - scope-26
|      |
|      |Cast[chararray] - scope-30
|      |
|      |---Project[bytearray][3] - scope-29
|      |
|      |Cast[chararray] - scope-33
|      |
|      |---Project[bytearray][4] - scope-32
|
|---student:
Load(hdfs://localhost:9000/pig_data/student_data.txt:PigStorage('
, ')) - scope19
2015-10-05 11:32:43,682 [main]
INFO
org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MRCo
mpiler -
File concatenation threshold: 100 optimistic? false
2015-10-05 11:32:43,684 [main]
INFO
org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.Mult
iQueryOp timizer -

```

```

MR plan size before optimization: 1 2015-10-05 11:32:43,685
[main]
INFO
org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.
MultiQueryOp timizer - MR plan size after optimization: 1
#-----
# Map Reduce Plan
#-----
MapReduce node scope-37
Map Plan
student: Store(fakefile:org.apache.pig.builtin.PigStorage) -
scope-36
|
|---student: New For Each(false,false,false,false,false) [bag] -
scope-35
|   |
|   Cast[int] - scope-21
|   |
|   |---Project[bytearray][0] - scope-20
|   |
|   Cast[chararray] - scope-24
|   |
|   |---Project[bytearray][1] - scope-23
|   |
|   Cast[chararray] - scope-27
|   |
|   |---Project[bytearray][2] - scope-26
|   |
|   Cast[chararray] - scope-30
|   |
|   |---Project[bytearray][3] - scope-29
|   |
|   Cast[chararray] - scope-33
|   |
|   |---Project[bytearray][4] - scope-32
|   |
|---student:
Load(hdfs://localhost:9000/pig_data/student_data.txt:PigStorage('
, ')) - scope
19----- Global sort: false
-----

```

Illustrate Operator

The **illustrate** operator gives you the step-by-step execution of a sequence of statements.

Syntax

Given below is the syntax of the **illustrate** operator.

```
grunt> illustrate Relation_name;
```

let us illustrate the relation named student as shown below.

```
grunt> illustrate student;
```

Output

On executing the above statement, you will get the following output.

```
grunt> illustrate student;
```

```
INFO
org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.PigM
apOnly$M ap - Aliases
being processed per job phase (AliasName[line,offset]): M:
student[1,10] C:  R:
-----
|student | id:int | firstname:chararray | lastname:chararray |
phone:chararray | city:chararray |
-----
|          | 002      | siddarth              | Battacharya          |
9848022338 | Kolkata    |
-----
-----
```

Group Operator

The **GROUP** operator is used to group the data in one or more relations. It collects the data having the same key.

Syntax

Given below is the syntax of the **group** operator.

```
grunt> Group_data = GROUP Relation_name BY age;
```

let us group the records/tuples in the relation by age as shown below.

```
grunt> group_data = GROUP student_details by age;
```

Verification

Verify the relation **group_data** using the **DUMP** operator as shown below.

```
grunt> Dump group_data;
```

Output

Then you will get output displaying the contents of the relation named **group_data** as shown below. Here you can observe that the resulting schema has two columns –

- One is **age**, by which we have grouped the relation.
- The other is a **bag**, which contains the group of tuples, student records with the respective age.

```
(21,{(4,Preethi,Agarwal,21,9848022330,Pune),(1,Rajiv,Reddy,21,984
8022337,Hyderabad)})
(22,{(3,Rajesh,Khanna,22,9848022339,Delhi),(2,siddarth,Battacharya,22,9848022338,Kolkata)})
(23,{(6,Archana,Mishra,23,9848022335,Chennai),(5,Trupthi,Mohanthy,23,9848022336,Bhuvaneshwar)})
(24,{(8,Bharathi,Nambiayar,24,9848022333,Chennai),(7,Komal,Nayak,24,9848022334,trivendram)})
```

You can see the schema of the table after grouping the data using the **describe** command as shown below.

```
grunt> Describe group_data;
```

```
group_data: {group: int, student_details: {(id: int, firstname:
chararray,
                lastname: chararray, age: int, phone:
chararray, city: chararray)}}
```

In the same way, you can get the sample illustration of the schema using the **illustrate** command as shown below.

```
$ Illustrate group_data;
```

It will produce the following output –

```
-----
-----
|group_data|  group:int |
student_details:bag{:tuple(id:int,firstname:chararray,lastname:ch
ararray,age:int,phone:chararray,city:chararray)}|
-----
-----
|          |      21      | { 4, Preethi, Agarwal, 21, 9848022330,
Pune), (1, Rajiv, Reddy, 21, 9848022337, Hyderabad)}|
|          |      2      | { (2, siddarth, Battacharya, 22, 9848022338, Kolkata), (003, Rajesh, Khann
a, 22, 9848022339, Delhi)}|
-----
-----
```

Grouping by Multiple Columns

Let us group the relation by age and city as shown below.

```
grunt> group_multiple = GROUP student_details by (age, city);
```

You can verify the content of the relation named **group_multiple** using the Dump operator as shown below.

```
grunt> Dump group_multiple;
```

```
((21, Pune), {(4, Preethi, Agarwal, 21, 9848022330, Pune)})
((21, Hyderabad), {(1, Rajiv, Reddy, 21, 9848022337, Hyderabad)})
((22, Delhi), {(3, Rajesh, Khanna, 22, 9848022339, Delhi)})
((22, Kolkata), {(2, siddarth, Battacharya, 22, 9848022338, Kolkata)})
((23, Chennai), {(6, Archana, Mishra, 23, 9848022335, Chennai)})
((23, Bhuwaneshwar), {(5, Trupthi, Mohanthy, 23, 9848022336, Bhuwaneshwa
r)})
((24, Chennai), {(8, Bharathi, Nambiayar, 24, 9848022333, Chennai)})
((24, trivendram), {(7, Komal, Nayak, 24, 9848022334, trivendram)})
```

Group All

You can group a relation by all the columns as shown below.

```
grunt> group_all = GROUP student_details All;
```

Now, verify the content of the relation **group_all** as shown below.

```
grunt> Dump group_all;
```



```
(all, {(8,Bharathi,Nambiayar,24,9848022333,Chennai),(7,Komal,Nayak,24,9848022334,Trivendram),
(6,Archana,Mishra,23,9848022335,Chennai),(5,Trupthi,Mohanthi,23,9848022336,Bhuwaneshwar),
(4,Preethi,Agarwal,21,9848022330,Pune),(3,Rajesh,Khanna,22,9848022339,Delhi),
(2,siddarth,Battacharya,22,9848022338,Kolkata),(1,Rajiv,Reddy,21,9848022337,Hyderabad)})
```

Cogroup Operator

The **COGROUP** operator works more or less in the same way as the **GROUP** operator. The only difference between the two operators is that the **group** operator is normally used with one relation, while the **cogroup** operator is used in statements involving two or more relations.

Grouping Two Relations using Cogroup

Assume that we have two files namely **student_details.txt** and **employee_details.txt** in the directory **/pig_data/** HDFS

employee_details.txt

```
001,Robin,22,newyork
002,BOB,23,Kolkata
003,Maya,23,Tokyo
004,Sara,25,London
005,David,23,Bhuwaneshwar
006,Maggy,22,Chennai
```

And we have loaded these files into Pig with the relation names **student_details** and **employee_details** respectively

Now, let us group the records/tuples of the relations **student_details** and **employee_details** with the key age, as shown below.

```
grunt> cogroup_data = COGROUP student_details by age,
employee_details by age;
```

Verification

Verify the relation **cogroup_data** using the **DUMP** operator as shown below.

```
grunt> Dump cogroup_data;
```

Output

It will produce the following output, displaying the contents of the relation named **cogroup_data** as shown below.

```
(21,{(4,Preethi,Agarwal,21,9848022330,Pune),
(1,Rajiv,Reddy,21,9848022337,Hyderabad)},
{ })
(22,{(3,Rajesh,Khanna,22,9848022339,Delhi),
(2,siddarth,Battacharya,22,9848022338,Kolkata)},
{(6,Maggy,22,Chennai),(1,Robin,22,newyork)})
(23,{(6,Archana,Mishra,23,9848022335,Chennai),(5,Trupthi,Mohanthi,23,9848022336,Bhuwaneshwar)},
{ })
```

```
{ (5, David, 23, Bhuwaneshwar), (3, Maya, 23, Tokyo), (2, BOB, 23, Kolkata) })
(24, { (8, Bharathi, Nambiayar, 24, 9848022333, Chennai), (7, Komal, Nayak,
24, 9848022334, trivendram) },
{ })
(25, { },
{ (4, Sara, 25, London) })
```

The **cogroup** operator groups the tuples from each relation according to age where each group depicts a particular age value.

For example, if we consider the 1st tuple of the result, it is grouped by age 21. And it contains two bags –

- the first bag holds all the tuples from the first relation (**student_details** in this case) having age 21, and
- the second bag contains all the tuples from the second relation (**employee_details** in this case) having age 21.

In case a relation doesn't have tuples having the age value 21, it returns an empty bag.