

19162121031

SMIT R PATEL

SME 5

HBASE PRACTICAL

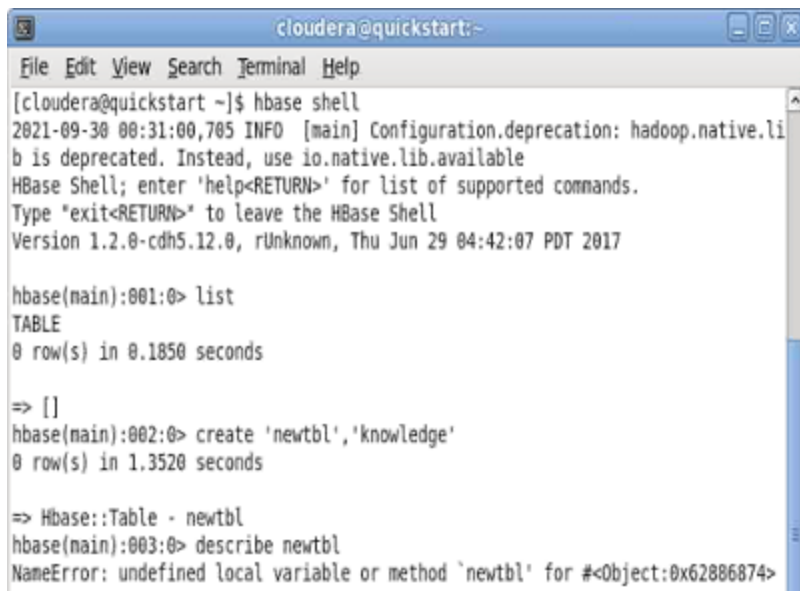
BDA

AIM- To understand basic commands in HBase.

Exercise: In order to work with an 'unstructured data' containing file, your sibling wants you to help her work in Hadoop. She realises that she might need to modify data and have random reads/writes, which traditional HDFS would not help with. Give her a solution and provide a way for her to work through this problem while teaching her how to work with that tool.

Tasks:-

1. Open the HBase shell and create a table with a column family.



```
cloudera@quickstart:~  
File Edit View Search Terminal Help  
[cloudera@quickstart ~]$ hbase shell  
2021-09-30 00:31:00,705 INFO [main] Configuration.deprecation: hadoop.native.lib  
is deprecated. Instead, use io.native.lib.available  
HBase Shell; enter 'help<RETURN>' for list of supported commands.  
Type "exit<RETURN>" to leave the HBase Shell  
Version 1.2.0-cdh5.12.0, rUnknown, Thu Jun 29 04:42:07 PDT 2017  
  
hbase(main):001:0> list  
TABLE  
0 row(s) in 0.1850 seconds  
  
=> []  
hbase(main):002:0> create 'newtbl','knowledge'  
0 row(s) in 1.3520 seconds  
  
=> Hbase::Table - newtbl  
hbase(main):003:0> describe newtbl  
NameError: undefined local variable or method `newtbl' for #<Object:0x62886874>
```

2. Work inside that column family and create columns with values.

```
[cloudera@quickstart Desktop]$ hbase shell
2021-09-30 00:45:33,239 INFO [main] Configuration.deprecation: hadoop.native.lib is deprecated. Instead, use io.native.lib.available
HBase Shell; enter 'help<RETURN>' for list of supported commands.
Type "exit<RETURN>" to leave the HBase Shell
Version 1.2.0-cdh5.12.0, rUnknown, Thu Jun 29 04:42:07 PDT 2017

hbase(main):001:0> status 'summary'
1 active master, 0 backup masters, 1 servers, 0 dead, 3.0000 average load

hbase(main):002:0> list
TABLE
newtbl
1 row(s) in 0.0340 seconds

=> ["newtbl"]
hbase(main):003:0> put 'newtbl','r1','knowledge:sports','chess'
0 row(s) in 0.1510 seconds

hbase(main):004:0> put 'newtbl','r1','knowledge:science','Maths'
0 row(s) in 0.0090 seconds

hbase(main):005:0> put 'newtbl','r1','knowledge:Music','Spiritual'
0 row(s) in 0.0140 seconds

hbase(main):006:0> put 'newtbl','r1','knowledge:Economics','Macro'
0 row(s) in 0.0050 seconds

hbase(main):007:0> put 'newtbl','r2','knowledge:Economics','Micro'
0 row(s) in 0.0060 seconds

hbase(main):008:0> put 'newtbl','r2','knowledge:Music','sleep'
0 row(s) in 0.0080 seconds

hbase(main):009:0> scan 'newtbl'
```

3. Query data from the above-mentioned table.

```
hbase(main):017:0> get 'newtbl', 'r1'
COLUMN                                CELL
knowledge:Economics                  timestamp=1632988279919, value=Macro
knowledge:Music                      timestamp=1632988245409, value=Spiritual
knowledge:science                    timestamp=1632988230605, value=Maths
knowledge:sports                     timestamp=1632988177749, value=chess
4 row(s) in 0.0320 seconds

hbase(main):018:0> get 'newtbl', 'r2'
COLUMN                                CELL
knowledge:Economics                  timestamp=1632988309609, value=Micro
knowledge:Music                      timestamp=1632988326358, value=sleep
2 row(s) in 0.0160 seconds

hbase(main):019:0> list
TABLE
newtbl
1 row(s) in 0.0100 seconds

=> ["newtbl"]
```

4. Give commands for enabling and disabling the table, and also alter data values inside it.

```
File Edit View Search Terminal Help
hbase(main):009:0> scan 'newtbl'
ROW          COLUMN+CELL
r1           column=knowledge:Economics, timestamp=1632988279919, value=Macro
r1           column=knowledge:Music, timestamp=1632988245409, value=Spiritual
r1           column=knowledge:science, timestamp=1632988230605, value=Maths
r1           column=knowledge:sports, timestamp=1632988177749, value=chess
r2           column=knowledge:Economics, timestamp=1632988309609, value=Micro
r2           column=knowledge:Music, timestamp=1632988326358, value=sleep
2 row(s) in 0.0750 seconds

hbase(main):010:0> is_enabled 'newtbl'
true
0 row(s) in 0.0380 seconds

hbase(main):011:0> disable 'newtbl'
0 row(s) in 2.3270 seconds
```

```
hbase(main):012:0> scan 'newtbl'  
ROW                COLUMN+CELL
```

ERROR: newtbl is disabled.

Scan a table; pass table name and optionally a dictionary of scanner specifications. Scanner specifications may include one or more of: TIMERANGE, FILTER, LIMIT, STARTROW, STOPROW, ROWPREFIXFILTER, TIMESTAMP, MAXLENGTH or COLUMNS, CACHE or RAW, VERSIONS, ALL_METRICS or METRICS

If no columns are specified, all columns will be scanned.

To scan all members of a column family, leave the qualifier empty as in 'col_family'.

The filter can be specified in two ways:

1. Using a filterString - more information on this is available in the Filter Language document attached to the HBASE-4176 JIRA
2. Using the entire package name of the filter.

If you wish to see metrics regarding the execution of the scan, the ALL_METRICS boolean should be set to true. Alternatively, if you would prefer to see only a subset of the metrics, the METRICS array can be defined to include the names of only the metrics you care about.

Some examples:

```
hbase> scan 'hbase:meta'  
hbase> scan 'hbase:meta', {COLUMNS => 'info:regioninfo'}  
hbase> scan 'ns1:t1', {COLUMNS => ['c1', 'c2'], LIMIT => 10, STARTROW => 'xyz'}  
hbase> scan 't1', {COLUMNS => ['c1', 'c2'], LIMIT => 10, STARTROW => 'xyz'}  
hbase> scan 't1', {COLUMNS => 'c1', TIMERANGE => [1303668804, 1303668904]}  
hbase> scan 't1', {REVERSED => true}  
hbase> scan 't1', {ALL_METRICS => true}  
hbase> scan 't1', {METRICS => ['RPC_RETRIES', 'ROWS_FILTERED']}  
hbase> scan 't1', {ROWPREFIXFILTER => 'row2', FILTER => "  
  (QualifierFilter (>=, 'binary:xyz')) AND (TimestampsFilter ( 123, 456))"}  
hbase> scan 't1', {FILTER =>  
  org.apache.hadoop.hbase.filter.ColumnPaginationFilter.new(1, 0)}  
hbase> scan 't1', {CONSISTENCY => 'TIMELINE'}
```

[Cloudera Live : Welco... [cloudera@quickstart:... [Master: quickstart.clo... cloudera@quickstart:~ cloudera@quickstart:~...

```
hbase(main):013:0> scan 'newtbl'alter 'newtbl', 'info'
SyntaxError: (hbase):13: syntax error, unexpected tIDENTIFIER

scan 'newtbl'alter 'newtbl', 'info'
      ^

hbase(main):014:0> alter 'newtbl', 'info'
Updating all regions with the new schema...
1/1 regions updated.
Done.
0 row(s) in 1.9910 seconds

hbase(main):015:0> enable 'newtbl'
0 row(s) in 1.3090 seconds

hbase(main):016:0> describe 'newtbl'
Table newtbl is ENABLED
newtbl
COLUMN FAMILIES DESCRIPTION
{NAME => 'info', DATA_BLOCK_ENCODING => 'NONE', BLOOMFILTER => 'ROW', REPLICATION_SCOPE => '0', COMPRESSION => 'NONE', VERSIONS => '1', TTL => 'FOREVER', MIN_VERSIONS => '0', KEEP_DELETED_CELLS => 'FALSE', BLOCKSIZE => '65536', IN_MEMORY => 'false', BLOCKCACHE => 'true'}
{NAME => 'knowledge', DATA_BLOCK_ENCODING => 'NONE', BLOOMFILTER => 'ROW', REPLICATION_SCOPE => '0', VERSIONS => '1', COMPRESSION => 'NONE', MIN_VERSIONS => '0', TTL => 'FOREVER', KEEP_DELETED_CELLS => 'FALSE', BLOCKSIZE => '65536', IN_MEMORY => 'false', BLOCKCACHE => 'true'}
2 row(s) in 0.0340 seconds
```

Extra :-

```
cloudera@quickstart:~/Desktop
File Edit View Search Terminal Help

hbase(main):023:0> is_enabled 'newtbl'
true
0 row(s) in 0.0330 seconds

hbase(main):024:0> scan 'newtbl'
ROW COLUMN+CELL
r1 column=knowledge:Economics, timestamp=1632988279919, value=Macro
r1 column=knowledge:Music, timestamp=1632988245409, value=Spiritual
r1 column=knowledge:science, timestamp=1632988230605, value=Maths
r1 column=knowledge:sports, timestamp=1632988177749, value=chess
r2 column=knowledge:Economics, timestamp=1632988309609, value=Micro
r2 column=knowledge:Music, timestamp=1632988326358, value=sleep
2 row(s) in 0.0470 seconds

hbase(main):025:0> enable 'newtbl'
0 row(s) in 0.0450 seconds

hbase(main):026:0> scan 'newtbl'
ROW COLUMN+CELL
r1 column=knowledge:Economics, timestamp=1632988279919, value=Macro
r1 column=knowledge:Music, timestamp=1632988245409, value=Spiritual
r1 column=knowledge:science, timestamp=1632988230605, value=Maths
r1 column=knowledge:sports, timestamp=1632988177749, value=chess
r2 column=knowledge:Economics, timestamp=1632988309609, value=Micro
r2 column=knowledge:Music, timestamp=1632988326358, value=sleep
2 row(s) in 0.0230 seconds
```



```
cloudera@quickstart: ~/Desktop
File Edit View Search Terminal Help

hbase(main):027:0> is_disabled 'newtbl'
false
0 row(s) in 0.0360 seconds

hbase(main):028:0> scan 'newtbl'
ROW COLUMN+CELL
r1 column=knowledge:Economics, timestamp=1632988279919, value=Macro
r1 column=knowledge:Music, timestamp=1632988245409, value=Spiritual
r1 column=knowledge:science, timestamp=1632988230605, value=Maths
r1 column=knowledge:sports, timestamp=1632988177749, value=chess
r2 column=knowledge:Economics, timestamp=1632988309609, value=Micro
r2 column=knowledge:Music, timestamp=1632988326358, value=sleep
2 row(s) in 0.0270 seconds

hbase(main):029:0> disable 'newtbl'
0 row(s) in 2.3000 seconds

hbase(main):030:0> scan 'newtbl'
ROW COLUMN+CELL

ERROR: newtbl is disabled.

Scan a table; pass table name and optionally a dictionary of scanner
specifications. Scanner specifications may include one or more of:
TIMERANGE, FILTER, LIMIT, STARTROW, STOPROW, ROWPREFIXFILTER, TIMESTAMP,
MAXLENGTH or COLUMNS, CACHE or RAW, VERSIONS, ALL_METRICS or METRICS

If no columns are specified, all columns will be scanned.
To scan all members of a column family, leave the qualifier empty as in
'col_family'.

The filter can be specified in two ways:
1. Using a filterString - more information on this is available in the
Filter Language document attached to the HBASE-4176 JIRA
2. Using the entire package name of the filter.

If you wish to see metrics regarding the execution of the scan, the
ALL_METRICS boolean should be set to true. Alternatively, if you would
```